

# Starting point

We would like to compare 2 different approaches to classify objects based on predictors.

Chosen approaches:

- kNN classifier
- Naive Bayes classifier

The comparison will be based on Accuracy, Recall, Precision and F1-Value of the models. The goal is to find the best model to predict "income" based on meaningful categorical features in the *census* data.

```
In [ ]: import pandas as pd
import numpy as np

from matplotlib import pyplot as plt
from sklearn.metrics import (
    confusion_matrix,
    classification_report,
    RocCurveDisplay,
    roc_curve,
    auc
)

from sklearn.model_selection import (
    GridSearchCV,
    train_test_split
)

from sklearn.naive_bayes import GaussianNB, CategoricalNB
```

## Data preparation

For both models the census.csv data will be used.

```
In [179.. dat = pd.read_csv("census.csv")
dat.head()
```

Out [179...

	age	workclass	fnlwgt	education	education.num	marital.status	occupation
0	39	State-gov	77516	Bachelors	13	Never-married	Adm-clerical
1	50	Self-emp-not-inc	83311	Bachelors	13	Married-civ-spouse	Exec-managerial
2	38	Private	215646	HS-grad	9	Divorced	Handlers-cleaners
3	53	Private	234721	11th	7	Married-civ-spouse	Handlers-cleaners
4	28	Private	338409	Bachelors	13	Married-civ-spouse	Prof-specialty

The target is **Income**

In [180... `dat[["income"]].value_counts()`

Out [180... income  
 <=50K 24720  
 >50K 7841  
 Name: count, dtype: int64

Our target variable doesn't show the balanced distribution within the groups. So we should care about creating stratified data samples during splitting on train-test samples.

We need to choose a few meaningful categorical features as predictors.

```
In [181... def analyze_categorical_columns(df):
    categorical_cols = df.select_dtypes(include=['object', 'category']).columns

    if len(categorical_cols) == 0:
        print("No categorical columns in DataFrame.")
        return

    for col in categorical_cols:
        print(f"Number of unique values: {df[col].nunique()}")
        print(f"{df[col].value_counts()}")
        print("-" * 30)

analyze_categorical_columns(dat)
```

Number of unique values: 9

workclass

Private	22696
---------	-------

Self-emp-not-inc	2541
------------------	------

Local-gov	2093
-----------	------

?	1836
---	------

State-gov	1298
-----------	------

Self-emp-inc	1116
--------------	------

Federal-gov	960
-------------	-----

Without-pay	14
-------------	----

Never-worked	7
--------------	---

Name: count, dtype: int64

Number of unique values: 16

education

HS-grad	10501
---------	-------

Some-college	7291
--------------	------

Bachelors	5355
-----------	------

Masters	1723
---------	------

Assoc-voc	1382
-----------	------

11th	1175
------	------

Assoc-acdm	1067
------------	------

10th	933
------	-----

7th-8th	646
---------	-----

Prof-school	576
-------------	-----

9th	514
-----	-----

12th	433
------	-----

Doctorate	413
-----------	-----

5th-6th	333
---------	-----

1st-4th	168
---------	-----

Preschool	51
-----------	----

Name: count, dtype: int64

Number of unique values: 7

marital.status

Married-civ-spouse	14976
--------------------	-------

Never-married	10683
---------------	-------

Divorced	4443
----------	------

Separated	1025
-----------	------

Widowed	993
---------	-----

Married-spouse-absent	418
-----------------------	-----

Married-AF-spouse	23
-------------------	----

Name: count, dtype: int64

Number of unique values: 15

occupation

Prof-specialty	4140
----------------	------

Craft-repair	4099
--------------	------

Exec-managerial	4066
-----------------	------

Adm-clerical	3770
--------------	------

Sales	3650
-------	------

Other-service	3295
---------------	------

Machine-op-inspct	2002
-------------------	------

?	1843
---	------

Transport-moving	1597
------------------	------

Handlers-cleaners	1370
-------------------	------

Farming-fishing	994
-----------------	-----

Tech-support	928
--------------	-----

Protective-serv	649
-----------------	-----

Priv-house-serv	149
-----------------	-----

Armed-Forces 9  
 Name: count, dtype: int64

-----  
 Number of unique values: 6

relationship  
 Husband 13193  
 Not-in-family 8305  
 Own-child 5068  
 Unmarried 3446  
 Wife 1568  
 Other-relative 981  
 Name: count, dtype: int64

-----  
 Number of unique values: 5

race  
 White 27816  
 Black 3124  
 Asian-Pac-Islander 1039  
 Amer-Indian-Eskimo 311  
 Other 271  
 Name: count, dtype: int64

-----  
 Number of unique values: 2

sex  
 Male 21790  
 Female 10771  
 Name: count, dtype: int64

-----  
 Number of unique values: 42

native.country  
 United-States 29170  
 Mexico 643  
 ? 583  
 Philippines 198  
 Germany 137  
 Canada 121  
 Puerto-Rico 114  
 El-Salvador 106  
 India 100  
 Cuba 95  
 England 90  
 Jamaica 81  
 South 80  
 China 75  
 Italy 73  
 Dominican-Republic 70  
 Vietnam 67  
 Guatemala 64  
 Japan 62  
 Poland 60  
 Columbia 59  
 Taiwan 51  
 Haiti 44  
 Iran 43  
 Portugal 37  
 Nicaragua 34  
 Peru 31  
 France 29  
 Greece 29  
 Ecuador 28

```

Ireland      24
Hong         20
Cambodia     19
Trinidad&Tobago 19
Laos         18
Thailand      18
Yugoslavia   16
Outlying-US(Guam-USVI-etc) 14
Honduras     13
Hungary      13
Scotland     12
Holand-Netherlands 1
Name: count, dtype: int64

```

```

-----
Number of unique values: 2
income
<=50K      24720
>50K        7841
Name: count, dtype: int64
-----

```

Based on information above the set of these categorical feachures were selected:

- workclass (9 classes)
- education (16 classes)
- marital.status (7 classes)
- occupation (15 classes)
- sex (2 classes)

```
In [182... whole_dataset = dat[["workclass","education","marital.status", "occupatio
whole_dataset.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 32561 entries, 0 to 32560
Data columns (total 6 columns):
#   Column          Non-Null Count  Dtype
---  ---
0   workclass       32561 non-null  object
1   education       32561 non-null  object
2   marital.status  32561 non-null  object
3   occupation      32561 non-null  object
4   sex             32561 non-null  object
5   income          32561 non-null  object
dtypes: object(6)
memory usage: 1.5+ MB

```

```
In [183... whole_dataset.isna().sum()
```

```

Out[183... workclass      0
education    0
marital.status  0
occupation   0
sex          0
income       0
dtype: int64

```

There are no any missing values, good.

## Onehot encoding

```
In [184... # Use pd.get_dummies() to one-hot encode the categorical columns
ds_encoded = pd.get_dummies(whole_dataset, columns=["workclass","educatio
ds_encoded.iloc[:,5].head(1)
```

```
Out [184...
income  workclass_Federal-  workclass_Local-  workclass_Never-  workclass_Pri
         gov                gov                worked
0  <=50K                False                False                False                F
```

Here we can see just first few columns.

```
In [185... print(f"Total number of columns: {len(ds_encoded.columns)}")
print(ds_encoded.columns)
```

```
Total number of columns: 45
Index(['income', 'workclass_Federal-gov', 'workclass_Local-gov',
      'workclass_Never-worked', 'workclass_Private', 'workclass_Self-emp-
inc',
      'workclass_Self-emp-not-inc', 'workclass_State-gov',
      'workclass_Without-pay', 'education_11th', 'education_12th',
      'education_1st-4th', 'education_5th-6th', 'education_7th-8th',
      'education_9th', 'education_Assoc-acdm', 'education_Assoc-voc',
      'education_Bachelors', 'education_Doctorate', 'education_HS-grad',
      'education_Masters', 'education_Preschool', 'education_Prof-schoo
l',
      'education_Some-college', 'marital.status_Married-AF-spouse',
      'marital.status_Married-civ-spouse',
      'marital.status_Married-spouse-absent', 'marital.status_Never-marri
ed',
      'marital.status_Separated', 'marital.status_Widowed',
      'occupation_Adm-clerical', 'occupation_Armed-Forces',
      'occupation_Craft-repair', 'occupation_Exec-managerial',
      'occupation_Farming-fishing', 'occupation_Handlers-cleaners',
      'occupation_Machine-op-inspct', 'occupation_Other-service',
      'occupation_Priv-house-serv', 'occupation_Prof-specialty',
      'occupation_Protective-serv', 'occupation_Sales',
      'occupation_Tech-support', 'occupation_Transport-moving', 'sex_Mal
e'],
      dtype='object')
```

The total number of columns is 45.

Now the data is ready to be splitted on test-train samples.

## Train-test split

```
In [186... X = ds_encoded.iloc[:,1:]
y = ds_encoded.iloc[:,0]

X_train, X_test, y_train, y_test = train_test_split(
```

```
X, y, test_size=1/3, random_state=471
)
```

```
In [187... pd.concat({"train": y_train.value_counts(), "test": y_test.value_counts()
# y_test.value_counts().merge(y_train.value_counts(), how = "inner")
```

```
Out [187...          train  test
```

**income**

	train	test
<=50K	16518	8202
>50K	5189	2652

Looks quite stratified.

## k-NN-model

As in previous homework we can use grid search to find good k and determine the best kNN model. We already know, that this method perform pretty well because it uses cross validation (in our case with 10 folds).

```
In [188... gs = GridSearchCV(estimator = knn,
                    param_grid = {'n_neighbors' : list(range(1,10))},
                    scoring = 'accuracy',
                    cv = 10,
                    refit = True)
best_knn_model = gs.fit(X_train, y_train).best_estimator_
print('Best k : %d' % best_knn_model.get_params()['n_neighbors'])
```

Best k : 9

Grid search tells us that the best k is 9.

We can calculate some performance metrics to check how good the model actually.

```
In [189... pred = best_knn_model.predict(X_test)
print(confusion_matrix(y_test, pred))
```

```
[[7442  760]
 [1284 1368]]
```

There are 7442 and 1368 correctly detected "<=50K" and ">50K" classes respectively.

But the number of false predictions is ratherly big (760 and 1284). Seems like this kNN model struggle to determine rich class (>50K), because the second row in the confusion matrix represented by numbers close to 1300, whereas the total number of observations with ">50K" in test data is around 2600. That means just half of the rich class was detected.

We can explore classification\_report to get more informations about model.

```
In [190... print(classification_report(y_test, pred))
```

	precision	recall	f1-score	support
<=50K	0.85	0.91	0.88	8202
>50K	0.64	0.52	0.57	2652
accuracy			0.81	10854
macro avg	0.75	0.71	0.73	10854
weighted avg	0.80	0.81	0.80	10854

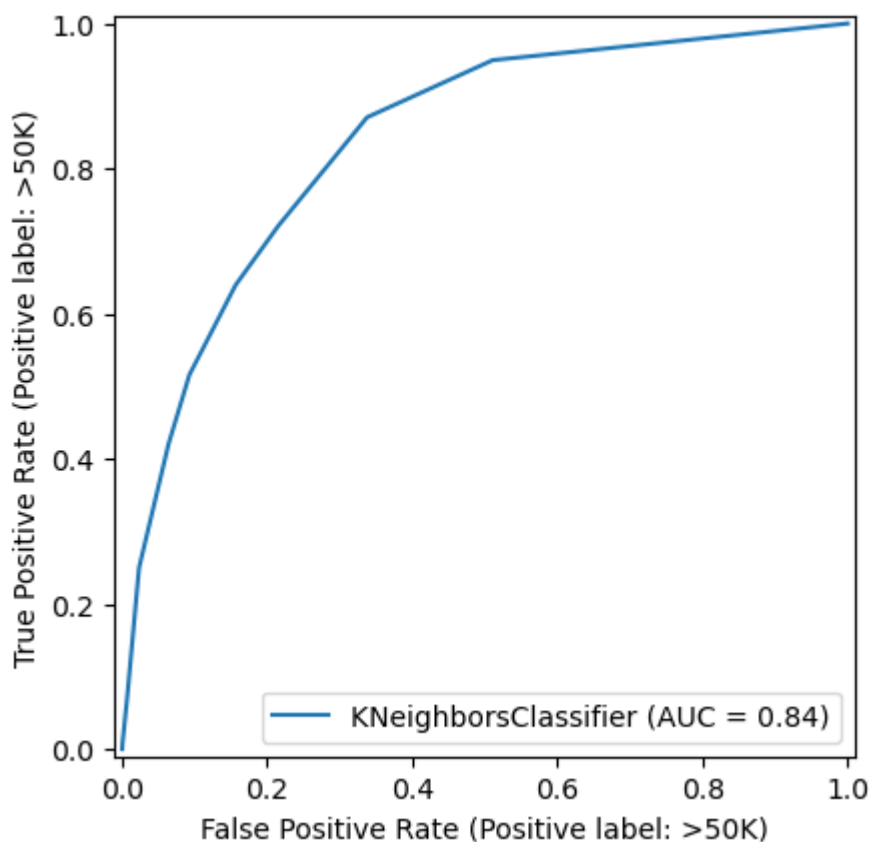
As mentioned before the class marked as ">50K" is problematic for the model, thus recall is just 0.52. Also the precision for this category is rather small showing 0.64.

The class with larger number of observations shows better performance with respect to precision (0.85) and recall (0.91), and thus the f1-score (0.88) as well.

The accuracy is showing 0.81, but such a big number reaches because of nice performance of larger class with about 80% of observations (<=50K), while smaller class (>50K) shows bad performance.

Combined performance metrics (macro avg & weighted avg) showcase around 0.73 and 0.80 respectively, which is not so good.

```
In [191]: RocCurveDisplay.from_estimator(best_knn_model, X_test, y_test)
plt.show()
```



## NaiveBayes model

Naive Bayes classifier assumes that the effect of a particular feature in a class is independent of other features



There are 2 types of Naive Bayes functions that exists in sklearn.

### GaussianNB:

- Assumes continuous features follow a Gaussian (normal) distribution.
- Used for classification with continuous features.

### CategoricalNB:

- Assumes categorical features follow a categorical distribution.
- Used for classification with categorical features.

In our case, we only have categorical features, so it's an easy decision to choose the latter method. However, it's still interesting to see how the former method might perform in this kind of task, because there was no *CategoricalNB* in previous versions of Scikit-learn. Additionally, I would like to test it because **ChatGPT** suggested using Gaussian NB, which seems strange when we only have categorical variables.

## Gaussian Naive Bayes

```
In [192... modelNB1 = GaussianNB()

modelNB1.fit(X_train, y_train)

predNB1 = modelNB1.predict(X_test)
```

```
In [193... print(confusion_matrix(y_test, predNB1))

[[1966 6236]
 [ 91 2561]]
```

Looks very bad, as expected because of use wrong Naive Bayes method (Gaussian instead of Categorical)

```
In [194... print(classification_report(y_test, predNB1))
```

	precision	recall	f1-score	support
<=50K	0.96	0.24	0.38	8202
>50K	0.29	0.97	0.45	2652
accuracy			0.42	10854
macro avg	0.62	0.60	0.42	10854
weighted avg	0.79	0.42	0.40	10854

Classification report also proof that this model performs bad with *weghted avg f1-score* around 0.4.

## Categorical Naive Bayes

```
In [195... modelNB2 = CategoricalNB()
```

```
modelNB2.fit(X_train, y_train)

predNB2 = modelNB2.predict(X_test)
```

```
In [196... print(confusion_matrix(y_test, predNB2))
```

```
[[6979 1223]
 [ 931 1721]]
```

This matrix looks much better compared to modelNB1.

But we would like to compare performance metrics with knn\_model.

```
In [197... print(classification_report(y_test, predNB2))
```

	precision	recall	f1-score	support
<=50K	0.88	0.85	0.87	8202
>50K	0.58	0.65	0.62	2652
accuracy			0.80	10854
macro avg	0.73	0.75	0.74	10854
weighted avg	0.81	0.80	0.80	10854

Usage of Categorical Naive Bayes brings us to model, that show simmlar patterns as knn\_model, but slightly different.

As well as knn\_model this NB\_model shows better performance regarding larger class (<=50K) with precision at 0.88 and recall at 0.85. In contrast to knn\_model the later metric is lower, but not significantly (0.91 -> 0.85 still high compared to other numbers).

As for smalle class (>50K) this model performs in another way, because the precision (0.64 -> 0.58) and recall (0.52 -> 0.65) swapped their positions in relative numbers. The f1-score for the larger group sligthly decreased (0.88 -> 0.87), while this for smaller group increased (0.57 -> 0.62).

To understand which model is better knn\_model or NB\_model, we should define what is more important for us: precision or recall.

In my personal oppinion the recall of samller class should be more important in this case, **therefore NB\_model is slightly better for this task**, while still the result of this model is 80% in terms of accuracy.

## Threshold tuning

```
In [198... X_train_part, X_val, y_train_part, y_val = train_test_split(
    X_train, y_train, test_size=1/3, random_state=1)
```

We splitted the train dataset, that we used befor, on 2 new datasets: train\_part and val.

Actual parts of the whole dataset:

- 1/3 test data
- 4/9 train\_part data
- 2/9 val data


```
In [199... modelNB3 = CategoricalNB()

modelNB3.fit(X_train_part, y_train_part)

y_prob_val = modelNB3.predict_proba(X_val)[:, 1]
fpr, tpr, thresholds = roc_curve(y_val, y_prob_val, pos_label='>50K')

roc_auc = auc(fpr, tpr)

plt.plot(fpr, tpr, label=f'ROC curve (area = {roc_auc:.2f})')
plt.plot([0, 1], [0, 1], 'k--')
plt.legend()
plt.show()
```



Another method to plot Receiver Operating Characteristic (ROC) curve using manually calculated FalsePositiveRate (fpr) and TruePositiveRate (tpr).

To define the best threshold we can use Youden's J Statistic, which is simply difference between *tpr* and *fpr*. That means that we try to find threshold that helps to get the model with better recall of the chosen class (in this case ">50K").

```
In [200... youden_j = tpr - fpr

best_threshold = thresholds[np.argmax(youden_j)]

y_prob_test = modelNB3.predict_proba(X_test)[:, 1]
pred_threshold = (y_prob_test >= best_threshold).astype(int)
```

```
print("threshold:", best_threshold)
prediction = ['>50K' if p == 1 else '<=50K' for p in pred_threshold]
```

threshold: 0.2461130910160678

New threshold is 0.246, which is significantly lower than basic one (0.5). We can assume the model will pick ">50K" more often, because the threshold becomes lower.

```
In [201... print(confusion_matrix(y_test, prediction))
```

```
[[6057 2145]
 [ 477 2175]]
```

As mentioned before the model increased predictions of ">50K" in total, bringing us to more false positive predictions (1223 -> 2145), while the number of true positive predictions increases slightly (1721 -> 2175).

TP and FP are almost equal, which means that the model produces 50% incorrect predictions for small class.

```
In [202... print(classification_report(y_test, prediction))
```

	precision	recall	f1-score	support
<=50K	0.93	0.74	0.82	8202
>50K	0.50	0.82	0.62	2652
accuracy			0.76	10854
macro avg	0.72	0.78	0.72	10854
weighted avg	0.82	0.76	0.77	10854

As expected, the recall for small class becomes higher, but we lose in precision: 50% is low.

Interesting thing that the f1-score for ">50K" doesn't change (0.62 -> 0.62).

Also, accuracy is slightly decreased, compared to model with initial threshold. As well as overall performance metrics becomes lower.

## Conclusion

Overall, all the models that we tested for this task can be reasonably used, except the model based on GaussianNB. Some of them have slightly better recall of small class, while others perform better in terms of precision for that. Also, the kNN model demonstrates that it requires much more computational power than NaiveBayes in this case, because we have a lot of dimensions (>40) after one-hot encoding.

I would say that the final model should be ClassificationNaiveBayes based on initial (0.5) threshold, because it shows balanced performance.