

# Random Forests

Vitalii Novikov

## Table of contents

<b>Libraries</b>	<b>1</b>
<b>Starting point</b>	<b>2</b>
<b>Data management</b>	<b>2</b>
Pick useful predictors . . . . .	4
Stratified data set . . . . .	5
Using 'sample' function . . . . .	5
Using caret library . . . . .	6
Normalize the data . . . . .	7
Scaled train dataset . . . . .	8
Scaled test dataset . . . . .	9
<b>Modeling part</b>	<b>9</b>
First Random Forest . . . . .	9
Hyperparameters tuning . . . . .	10
Final Model . . . . .	12
Possible improvements . . . . .	14

## Libraries

```
library <- function(...) {suppressPackageStartupMessages(base::library(...))}  
if (!require(caret)) install.packages("caret"); library(caret)
```

Loading required package: caret

Loading required package: ggplot2

Loading required package: lattice

```
if (!require(mlbench)) install.packages("mlbench"); library(mlbench)
```

Loading required package: mlbench

```
if (!require(dplyr)) install.packages("dplyr"); library(dplyr)
```

Loading required package: dplyr

Attaching package: 'dplyr'

The following objects are masked from 'package:stats':

filter, lag

The following objects are masked from 'package:base':

intersect, setdiff, setequal, union

```
library(plyr)
library(randomForest)
library(pROC)
```

## Starting point

We would like to use Random Forests to predict the diagnosis (Benign or Malignant based on 10 features: - radius (mean of distances from center to points on the perimeter) - texture (standard deviation of gray-scale values) - perimeter - area - smoothness (local variation in radius lengths) - compactness ( $\text{perimeter}^2 / \text{area} - 1.0$ ) - concavity (severity of concave portions of the contour) - concave points (number of concave portions of the contour) - symmetry - fractal dimension ("coastline approximation" - 1)

## Data management

```
dat <- read.csv("data.csv")
glimpse(dat)
```

Rows: 569

Columns: 33

```
$ id          <int> 842302, 842517, 84300903, 84348301, 84358402, ~
$ diagnosis   <chr> "M", "M", "M", "M", "M", "M", "M", "M", "M", "~
$ radius_mean <dbl> 17.990, 20.570, 19.690, 11.420, 20.290, 12.450~
$ texture_mean <dbl> 10.38, 17.77, 21.25, 20.38, 14.34, 15.70, 19.9~
$ perimeter_mean <dbl> 122.80, 132.90, 130.00, 77.58, 135.10, 82.57, ~
$ area_mean    <dbl> 1001.0, 1326.0, 1203.0, 386.1, 1297.0, 477.1, ~
$ smoothness_mean <dbl> 0.11840, 0.08474, 0.10960, 0.14250, 0.10030, 0~
$ compactness_mean <dbl> 0.27760, 0.07864, 0.15990, 0.28390, 0.13280, 0~
$ concavity_mean <dbl> 0.30010, 0.08690, 0.19740, 0.24140, 0.19800, 0~
$ concave.points_mean <dbl> 0.14710, 0.07017, 0.12790, 0.10520, 0.10430, 0~
$ symmetry_mean <dbl> 0.2419, 0.1812, 0.2069, 0.2597, 0.1809, 0.2087~
$ fractal_dimension_mean <dbl> 0.07871, 0.05667, 0.05999, 0.09744, 0.05883, 0~
$ radius_se     <dbl> 1.0950, 0.5435, 0.7456, 0.4956, 0.7572, 0.3345~
$ texture_se     <dbl> 0.9053, 0.7339, 0.7869, 1.1560, 0.7813, 0.8902~
$ perimeter_se   <dbl> 8.589, 3.398, 4.585, 3.445, 5.438, 2.217, 3.18~
$ area_se        <dbl> 153.40, 74.08, 94.03, 27.23, 94.44, 27.19, 53.~
$ smoothness_se  <dbl> 0.006399, 0.005225, 0.006150, 0.009110, 0.0114~
$ compactness_se <dbl> 0.049040, 0.013080, 0.040060, 0.074580, 0.0246~
$ concavity_se   <dbl> 0.05373, 0.01860, 0.03832, 0.05661, 0.05688, 0~
$ concave.points_se <dbl> 0.015870, 0.013400, 0.020580, 0.018670, 0.0188~
$ symmetry_se     <dbl> 0.03003, 0.01389, 0.02250, 0.05963, 0.01756, 0~
$ fractal_dimension_se <dbl> 0.006193, 0.003532, 0.004571, 0.009208, 0.0051~
$ radius_worst    <dbl> 25.38, 24.99, 23.57, 14.91, 22.54, 15.47, 22.8~
$ texture_worst    <dbl> 17.33, 23.41, 25.53, 26.50, 16.67, 23.75, 27.6~
$ perimeter_worst  <dbl> 184.60, 158.80, 152.50, 98.87, 152.20, 103.40,~
$ area_worst       <dbl> 2019.0, 1956.0, 1709.0, 567.7, 1575.0, 741.6, ~
$ smoothness_worst <dbl> 0.1622, 0.1238, 0.1444, 0.2098, 0.1374, 0.1791~
$ compactness_worst <dbl> 0.6656, 0.1866, 0.4245, 0.8663, 0.2050, 0.5249~
$ concavity_worst  <dbl> 0.71190, 0.24160, 0.45040, 0.68690, 0.40000, 0~
$ concave.points_worst <dbl> 0.26540, 0.18600, 0.24300, 0.25750, 0.16250, 0~
$ symmetry_worst   <dbl> 0.4601, 0.2750, 0.3613, 0.6638, 0.2364, 0.3985~
$ fractal_dimension_worst <dbl> 0.11890, 0.08902, 0.08758, 0.17300, 0.07678, 0~
$ X                <lgl> NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA~
```

```
nrow(dat)
```

```
[1] 569
```

This dataset include 569 observations. We will only use the “worst-features” for creating neural network.

## Pick useful predictors

```
clean_data <- as_tibble(dat[c(2,23:32)])
clean_data <- mutate(clean_data,
  diagnosis = case_match(diagnosis, "M" ~ "malignant", "B"~ "benign")
) %>% mutate(diagnosis = factor(diagnosis))
# %>% mutate(diagnosis = relevel(diagnosis, ref = "malignant"))
#clean_data[['diagnosis']] <- as.factor(clean_data[['diagnosis']])
clean_data |> head(2)
```

```
# A tibble: 2 x 11
  diagnosis radius_worst texture_worst perimeter_worst area_worst
<fct>      <dbl>      <dbl>      <dbl>      <dbl>
1 malignant    25.4        17.3        185.        2019
2 malignant    25.0        23.4        159.        1956
# i 6 more variables: smoothness_worst <dbl>, compactness_worst <dbl>,
#   concavity_worst <dbl>, concave.points_worst <dbl>, symmetry_worst <dbl>,
#   fractal_dimension_worst <dbl>
```

```
clean_data |> apply(2, function(x) sum(is.na(x)))
```

diagnosis	radius_worst	texture_worst
0	0	0
perimeter_worst	area_worst	smoothness_worst
0	0	0
compactness_worst	concavity_worst	concave.points_worst
0	0	0
symmetry_worst	fractal_dimension_worst	
0	0	

There is no missing data, good.


## Stratified data set

### Using 'sample' function

```
set.seed(230)
index <- sample(1:nrow(clean_data), round(0.8*nrow(clean_data)))
train_data <- clean_data[index,]
test_data <- clean_data[-index,]

pie_with_percentages <- function(diag_list, main = "Diagnosis", round = 1) {
  diag_table <- table(diag_list)
  diag_percentages <- prop.table(diag_table) * 100
  labels <- paste(names(diag_table), "\n", round(diag_percentages, round), "%")
  return(pie(diag_table, labels = labels, main = main))
}

layout(matrix(c(1,2,3), nrow = 1), respect = TRUE)
pie_with_percentages(train_data$diagnosis, main = "train_data")
pie_with_percentages(test_data$diagnosis, main = "test_data")
pie_with_percentages(clean_data$diagnosis, main = "original_data")
```



```
layout(1)
```


The test data show slightly lowest proportion of malignant then both train and original data, but it is still very close to actual picture.

### Using caret library

```
set.seed(123)
index_2 <- createDataPartition(clean_data$diagnosis, p = 0.8, list = FALSE)

train_data_2 <- clean_data[index_2,]
test_data_2 <- clean_data[-index_2,]

layout(matrix(c(1,2,3), nrow = 1), respect = TRUE)
pie_with_percentages(train_data_2$diagnosis, main = "train_data_2", round = 2)
pie_with_percentages(test_data_2$diagnosis, main = "test_data_2", round = 2)
pie_with_percentages(clean_data$diagnosis, main = "original_data", round = 2)
```



In this case we got really stratified train and test split of the data, because the proportions of classes are almost the same in each dataset. For further steps we will use this split of the data.

```

trainD <- train_data_2
testD <- test_data_2

split_description <- t(rbind(
  data.frame("trainD" = nrow(trainD), "testD" = nrow(testD)),
  c(round(nrow(trainD)/nrow(clean_data)*100,2),
    round(nrow(testD)/nrow(clean_data)*100,2))
))
colnames(split_description) <- c("nrow","percentage (%)")
split_description

```

	nrow	percentage (%)
trainD	456	80.14
testD	113	19.86

The trainD include 456 rows, which is 80.14% of the original dataset.

## Normalize the data

All predictors have their own scales. We should perform min-max-normalization.

```
trainD |> summary()
```

diagnosis	radius_worst	texture_worst	perimeter_worst
benign :286	Min. : 7.93	Min. :12.02	Min. : 50.41
malignant:170	1st Qu.:12.87	1st Qu.:21.05	1st Qu.: 83.73
	Median :14.91	Median :25.27	Median : 97.26
	Mean :16.19	Mean :25.65	Mean :106.72
	3rd Qu.:18.38	3rd Qu.:29.91	3rd Qu.:124.15
	Max. :36.04	Max. :47.16	Max. :251.20
area_worst	smoothness_worst	compactness_worst	concavity_worst
Min. : 185.2	Min. :0.07117	Min. :0.02729	Min. :0.0000
1st Qu.: 509.4	1st Qu.:0.11650	1st Qu.:0.14010	1st Qu.:0.1081
Median : 684.0	Median :0.13145	Median :0.21165	Median :0.2112
Mean : 875.0	Mean :0.13234	Mean :0.25314	Mean :0.2691
3rd Qu.:1035.0	3rd Qu.:0.14600	3rd Qu.:0.33145	3rd Qu.:0.3814
Max. :4254.0	Max. :0.21840	Max. :1.05800	Max. :1.1700
concave.points_worst	symmetry_worst	fractal_dimension_worst	
Min. :0.00000	Min. :0.1565	Min. :0.05504	
1st Qu.:0.06330	1st Qu.:0.2478	1st Qu.:0.07190	

Median :0.09766	Median :0.2813	Median :0.08009
Mean :0.11275	Mean :0.2881	Mean :0.08430
3rd Qu.:0.16025	3rd Qu.:0.3174	3rd Qu.:0.09192
Max. :0.29030	Max. :0.6638	Max. :0.20750

## Scaled train dataset

```
maxs <- apply(trainD[-1], 2, max)
mins <- apply(trainD[-1], 2, min)

scaled_trainD <- scale(trainD[-1], center = mins, scale = maxs - mins) %>%
  cbind(trainD[1])

summary(scaled_trainD)
```

radius_worst	texture_worst	perimeter_worst	area_worst
Min. :0.0000	Min. :0.0000	Min. :0.0000	Min. :0.00000
1st Qu.:0.1758	1st Qu.:0.2570	1st Qu.:0.1659	1st Qu.:0.07969
Median :0.2483	Median :0.3769	Median :0.2333	Median :0.12258
Mean :0.2937	Mean :0.3878	Mean :0.2804	Mean :0.16952
3rd Qu.:0.3716	3rd Qu.:0.5092	3rd Qu.:0.3672	3rd Qu.:0.20886
Max. :1.0000	Max. :1.0000	Max. :1.0000	Max. :1.00000
smoothness_worst	compactness_worst	concavity_worst	concave.points_worst
Min. :0.0000	Min. :0.0000	Min. :0.00000	Min. :0.0000
1st Qu.:0.3079	1st Qu.:0.1094	1st Qu.:0.09237	1st Qu.:0.2181
Median :0.4094	Median :0.1789	Median :0.18056	Median :0.3364
Mean :0.4155	Mean :0.2191	Mean :0.23002	Mean :0.3884
3rd Qu.:0.5083	3rd Qu.:0.2951	3rd Qu.:0.32598	3rd Qu.:0.5520
Max. :1.0000	Max. :1.0000	Max. :1.00000	Max. :1.0000
symmetry_worst	fractal_dimension_worst	diagnosis	
Min. :0.0000	Min. :0.0000	benign :286	
1st Qu.:0.1799	1st Qu.:0.1106	malignant:170	
Median :0.2461	Median :0.1643		
Mean :0.2595	Mean :0.1919		
3rd Qu.:0.3172	3rd Qu.:0.2419		
Max. :1.0000	Max. :1.0000		

Now we have all predictors in the same scale.



## Scaled test dataset

```
maxs <- apply(testD[-1], 2, max)
mins <- apply(testD[-1], 2, min)

scaled_testD <- scale(testD[-1], center = mins, scale = maxs - mins) %>%
  cbind(testD[1])
```

## Modeling part

### First Random Forest

The initial model include 20 trees and has 3 variables randomly sampled as candidates at each split ( $\text{floor}(\sqrt{10}) = 3$ ).

```
set.seed(333)
rf_model <- randomForest(diagnosis ~ .,
                        data = scaled_trainD,
                        ntree = 20,
                        mtry = floor(sqrt(10))
                        )
```

```
predictions <- predict(rf_model, newdata = scaled_testD)
prob_predictions <- predict(rf_model, newdata = scaled_testD, type = "prob")[,2]
print(confusionMatrix(predictions, scaled_testD$diagnosis, mode = "prec_recall"))
```

### Confusion Matrix and Statistics

	Reference	
Prediction	benign	malignant
benign	67	1
malignant	4	41

Accuracy : 0.9558  
95% CI : (0.8998, 0.9855)  
No Information Rate : 0.6283  
P-Value [Acc > NIR] : <2e-16

Kappa : 0.9066

McNemar's Test P-Value : 0.3711

Precision : 0.9853  
Recall : 0.9437  
F1 : 0.9640  
Prevalence : 0.6283  
Detection Rate : 0.5929  
Detection Prevalence : 0.6018  
Balanced Accuracy : 0.9599

'Positive' Class : benign

The initial model performed not so bad, just 5 observations were wrong classified. Therefore the Accuracy is about 96%.

## Hyperparameters tuning

As Hyperparameters we have: 1) number of trees in the forest 2) number of variables randomly sampled as candidates at each split (it is recommended to take around  $\sqrt{N\_predictors}$ ), but it is interesting to explore how it affect the model)

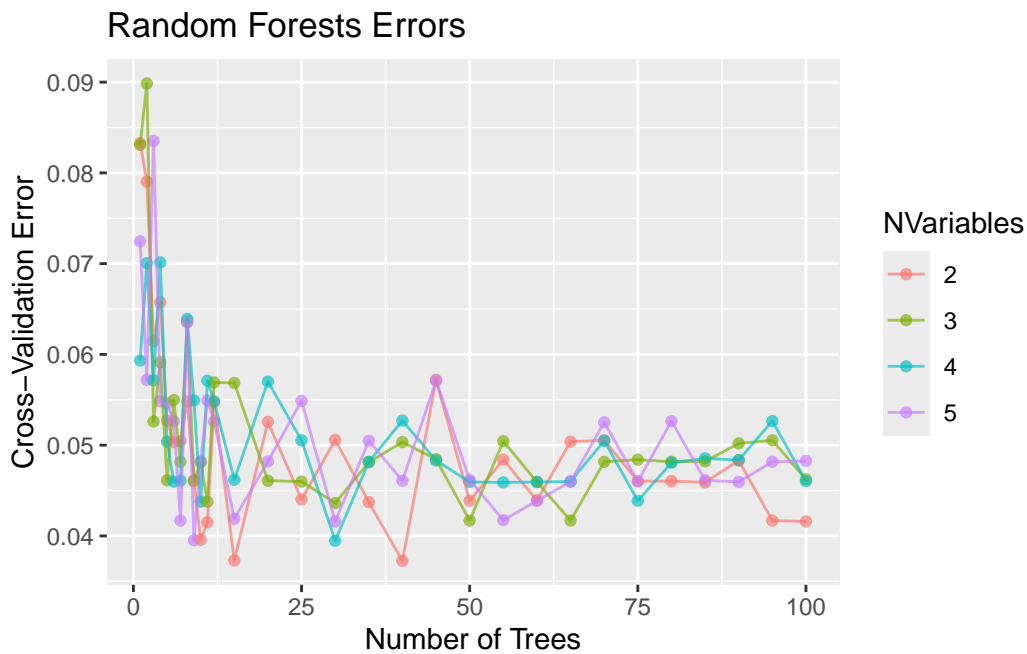
```
set.seed(1)
n_trees <- c(1:12, seq(15, 100, by = 5))
error_df <- data.frame(Trees = numeric(), NVariables = factor(), Error = numeric())

control <- trainControl(method = "cv", number = 10, classProbs = TRUE, savePredictions = "final")

for (ntree in n_trees) {
  for (nv_count in 2:5) {## number of variables randomly sampled as candidates at each split
    rf_model <- train(diagnosis ~ ., data = scaled_trainD,
                      method = "rf",
                      trControl = control,
                      tuneGrid = data.frame(mtry = nv_count),
                      ntree = ntree)
    # Collect cross-validated errors
    error_df <- rbind(error_df,
                      data.frame(Trees = ntree,
                                NVariables = as.character(nv_count),
                                Error = 1 - max(rf_model$results$Accuracy)))
  }
}
```

```
}}
```

```
ggplot(error_df, aes(x = Trees, y = Error, color = NVariables)) +
  geom_point(alpha = 0.65) +
  geom_line(alpha = 0.65) +
  #scale_x_log10()+
  labs(title = "Random Forests Errors",
        x = "Number of Trees",
        y = "Cross-Validation Error")
```




There is no special patterns in differences between number of variables sampled as candidates at each split. May be “3” is a little bit more stable, but still the errors line fluctuates between 0.04 and 0.05.

To easier interpret the start of the plateau we should use log-transformed x-axis.

```
ggplot(error_df, aes(x = Trees, y = Error, color = NVariables)) +
  geom_point(alpha = 0.65) +
  geom_line(alpha = 0.65) +
  scale_x_log10()+
  labs(title = "Random Forests Errors (log transformed)",
```

```
x = "Number of Trees",
y = "Cross-Validation Error")
```



The plateau begins around the number of trees equals 5, but there is a strange increase of errors for all models when number of trees equals 8.

When the model have 5 trees, the lowest error has the model with number of sampled variables = 3. Therefore, this model will be used in further steps.

## Final Model

As described before, cross validation tell us that the appropriate number of trees is 5 and the appropriate number of sampled variables is 3.

```
set.seed(55)
final_model <- randomForest(diagnosis ~ .,
                             data = scaled_trainD,
                             ntree = 5,
                             mtry = 3)

predictions <- predict(final_model, newdata = scaled_testD)
prob_predictions <- predict(final_model, newdata = scaled_testD, type = "prob")[,2]
print(confusionMatrix(predictions, scaled_testD$diagnosis, mode = "prec_recall"))
```

## Confusion Matrix and Statistics

	Reference	
Prediction	benign	malignant
benign	70	2
malignant	1	40

Accuracy : 0.9735  
95% CI : (0.9244, 0.9945)  
No Information Rate : 0.6283  
P-Value [Acc > NIR] : <2e-16

Kappa : 0.9429

McNemar's Test P-Value : 1

Precision : 0.9722  
Recall : 0.9859  
F1 : 0.9790  
Prevalence : 0.6283  
Detection Rate : 0.6195  
Detection Prevalence : 0.6372  
Balanced Accuracy : 0.9691


'Positive' Class : benign

The final model performed not bad, just 3 observations were wrong classified. Therefore the Balanced Accuracy around 97%. According to confusion matrix in 2 cases the malignant were classified as benign, which is bad.

```
roc_obj <- roc(scaled_testD$diagnosis, prob_predictions,  
              levels = c("benign", "malignant"))
```

Setting direction: controls < cases

```
plot(roc_obj, col = "darkblue", main = "ROC Curve")  
text(0.2, 0.2, labels = paste0("AUC = ", round(auc(roc_obj), 3)), cex = 1.2)
```



**ROC Curve** looks almost perfect with  $AUC = 0.994$ . There could be some improvements, because the TPR (sensitivity) become 1 at about 0.8 of specificity.

### Possible improvements

The final model is already quite good in terms of detecting “malignant”. If we want to detect all of the malignant items we can change the threshold to achieve the appropriate results, but it can lead to increase of False Positive predictions, which sometimes is ok. Also, the possible solution could be the “unsure” status for some range of probabilities. But for better calculating of percentages it may be necessary to use larger number of trees in the random forest.

And of course increasing of the number of observations in the initial dataset can help to train better model.